

WHOOOP

# Ranking Complaint Themes by Churn & Revenue Exposure

Prepared for WHOOOP Business Intelligence & Analytics

# A complaint list is an analysis. Revenue exposure is a business case.

WHOOP's own investor narrative is retention. This project ranks its recurring product complaints by how much they threaten that narrative — not by how loud they are.

**\$10.1B**

Series G valuation, Mar 2026

**2.5M**

members

**+103%**

YoY bookings growth

**\$1.1B**

bookings run-rate

*So what: any complaint theme that plausibly erodes retention is a threat to the number leadership just sold to investors — that's the lens this project uses to prioritize.*

# The KPI framework

Four metrics, each building on the last, to separate how loud a complaint is from how much it threatens revenue.

01

## Volume share

% of total complaint volume a theme represents.

02

## Churn-language share

% of that theme's complaints containing explicit cancel / switch / refund language.

03

## Churn-risk index (0–100)

Volume share × churn-language share, re-normalized across all 9 themes. The core KPI.


04

## Illustrative revenue at risk

Churn-risk share × an at-risk revenue pool (2.5M members × a stated, tunable churn assumption × \$440 ACV).

# The loudest complaints aren't the ones driving churn

% of total complaint volume vs. churn-risk index, by theme




# Ranked findings

*Illustrative annual revenue at risk assumes a stated, tunable 3% complaint-driven churn rate across the 2.5M member base.*

| Theme                                                     | Volume % | Churn-risk | \$ at risk / yr |
|-----------------------------------------------------------|----------|------------|-----------------|
| Subscription, pricing, cancellation & support             | 22       | 51.2       | \$16.90M        |
| Hardware reliability, device failure & charging           | 14       | 22.5       | \$7.42M         |
| App experience: sync, crashes, redesign, platform split   | 12       | 8.6        | \$2.84M         |
| Recovery-score accuracy ("doesn't match how I feel")      | 20       | 5.7        | \$1.88M         |
| Strain score confusion & strength-training under-counting | 12       | 4.3        | \$1.42M         |
| Heart-rate accuracy, high-intensity/weightlifting         | 10       | 3.2        | \$1.06M         |
| Band/clasp durability & skin irritation                   | 4        | 2.1        | \$0.69M         |
| Blood Pressure Insights / FDA (MG-specific)               | 1        | 1.3        | \$0.43M         |
| Sleep-tracking accuracy                                   | 5        | 1.1        | \$0.36M         |

# Complaint intensity over time, keyed to events

Illustrative relative index (0–100), calibrated to the sourced pattern of press/FDA/forum activity



May 2025: 5.0/MG launch + upgrade-fee reversal (2,400-upvote cancel post) • Jul 2025: FDA warning letter • Jun 2026: FDA closeout

# Recommendations by team



## Subscription, billing & cancellation

Churn-risk 51.2 • ~\$16.9M/yr

Instrument the cancellation funnel; quantify AI-only support tickets; check product behavior against the auto-renewal class action's allegations.



## Device failure & charging

Churn-risk 22.5 • ~\$7.4M/yr

Cohort view of MG/5.0 units from first pairing: hours-to-first-failure, replacement success rate, retention impact of the free-upgrade policy.



## Recovery, strain & HR accuracy

~13 combined • ~37% of volume

Prioritize communication/UX over further model tuning repeated firmware fixes haven't eliminated the complaint, so the gap looks like expectation-setting.



## Blood Pressure Insights / FDA

Churn-risk 1.3 • highest severity

Standing trust/regulatory-exposure metric for every new clinically-framed feature, given the closed FDA letter and pending class action.

# How confident should you be in these numbers

## Real, sourced evidence

21 attributed, verbatim quotes from press, FDA/court filings, WHOOP's own forum & support docs, and review platforms.

A separate, live dataset of 152 real reviews and comments — App Store, Google Play, Trustpilot, and Reddit — was hand-collected and is computed directly in Amplitude (see final slide).

## Transparent, tunable model

Category volume/churn shares and the 3% churn-rate assumption are named, stated, and openly labeled as estimates — not raw scraped counts.

Every number is re-computable by changing one input, and the full logic is written as Snowflake/dbt SQL models.

---

## To operationalize this at WHOOP:

- ▶ Wire the dbt project's raw sources into Snowflake; replace estimates with measured volumes
- ▶ Swap the regex category tagger for an LLM classifier, using regex as an explainable QA baseline
- ▶ Join to real cancellation/billing events to replace the 3% assumption with measured attrition
- ▶ The self-serve version of this already exists, live in Amplitude on the real 152-review dataset (final slide)

# Next steps

**Stand up the churn-risk index** as a shared metric alongside NPS/CSAT prioritize roadmap by revenue exposure, not ticket volume

**Instrument the real cancellation funnel** replace the 3% illustrative churn assumption with measured, cohort-based attrition

**Extend the Advanced Labs & BPI scorecards** proposed in the Healthcare Product deep-dive, before the next clinically-framed feature launches

**Supporting materials** full case study and SQL models available on request; live Amplitude dashboard (152 real reviews):  
[app.amplitude.com/analytics/share/413bf738048f4484809478a1e5f894b3](https://app.amplitude.com/analytics/share/413bf738048f4484809478a1e5f894b3)